




Training: soft masks via Gumbel-Sigmoid

Inference: hard masks via TEAL global top-k