

Output-Distribution Distillation for Contextual Sparsity Prediction in SwiGLU LLMs

Anonymous ACL submission

Abstract

Learned sparsity predictors for SwiGLU LLMs predominantly use binary cross-entropy (BCE), training each neuron independently, yet inference allocates budgets globally via top- k ranking. We train lightweight predictors ($<1\%$ of base parameters) with forward KL divergence, coupling all mask decisions through the output distribution. On LLaMA-3.1-8B at 30% sparsity, the KL predictor reduces perplexity by 14.5% over TEAL at matched sparsity, with directional commonsense gains ($n=3$) but knowledge-intensive regression (8B MMLU -5.2 pp; Qwen-14B GSM8K -19.9 pp; directional at $n=3$); post-hoc compensation (+67.2M params) becomes redundant. Output-distribution optimization yields masks that diverge 46–68% from magnitude-based targets; a target-swap ablation suggests these discovered patterns carry most of the quality information: BCE trained on KL-derived targets approaches KL-level PPL (PPL gap ≤ 0.65 pp; 3-seed, 30%). Per-layer controls show the signal type gap (9.47 PPL) exceeds the coupling breadth range (2.08 PPL) at 30%, though the decomposition is confounded by train-inference mismatch; a $K=4$ control suggests task-dependent coupling effects (§4.5). This is primarily an empirical study; wall-clock speedup requires sparse kernel support not addressed here. Cross-architecture evaluation on Mistral-7B-v0.3 and Qwen-2.5-14B confirms PPL generalization with scale-dependent downstream effects.

1 Introduction

How the training objective shapes learned sparsity prediction quality is an open question for SwiGLU LLMs (Shazeer, 2020). SwiGLU FFN blocks exhibit high contextual sparsity, but their gated activation makes importance prediction challenging: neuron importance depends on both gate and up-projection values, a target that magnitude-based methods handle suboptimally. Learned sparsity

predictors (Liu et al., 2023; Gautam et al., 2026) exploit this property by training small networks to anticipate which neurons are important, but they predominantly rely on binary cross-entropy (BCE) loss, treating each neuron as an independent classification target. Yet at inference, sparsity budgets are allocated globally: the TEAL (training-free activation sparsity; Liu et al., 2025) top- k mechanism ranks all neurons across all layers and retains only those above a model-wide threshold. This mismatch between per-neuron training and global inference allocation has not been formally analyzed. Unlike MaskLLM (Fang et al., 2024), which applies KL to static N:M weight masks, we address contextual activation sparsity where masks must generalize to unseen inputs, vary per token, and coordinate across layers through global allocation.

Gradient analysis reveals a structural asymmetry: BCE’s gradient $\partial \mathcal{L}_{\text{BCE}} / \partial s_j = \sigma(s_j) - m_j^*$ depends only on neuron j , while under KL divergence, the gradient couples each neuron to all others through the downstream Jacobian (§3.6). KL’s coupled gradients discover neuron importance patterns that diverge 68% from magnitude-based masks (mean per-layer Jaccard distance). A target-swap ablation shows KL-discovered mask targets carry most of the quality information: BCE re-trained on KL-derived targets approaches KL PPL (≤ 0.65 pp gap, 3-seed), while magnitude-trained BCE yields 28.13 ± 0.20 . Per-layer controls suggest coupling effects are task-dependent, with larger impacts on commonsense than knowledge tasks (§4.5).

We instantiate a lightweight predictor (75.9M params, $<1\%$ of base) trained with forward KL and Gumbel-Sigmoid relaxation (Jang et al., 2017) on LLaMA-3.1-8B (Grattafiori et al., 2024). At 30% sparsity under uniform allocation, the predictor reduces TEAL’s perplexity by 14.5% at matched sparsity (8.81 vs. 10.30, both at 28.5%) with downstream (MMLU $40.7 \pm 1.3\%$ vs. 65.3% dense, 3-

seed); global top- k amplifies cross-seed variance ($7.3\times$). At 50%, the perplexity advantage widens (50% over TEAL), but downstream reasoning collapses regardless of method.

Cross-architecture evaluation on Mistral-7B-v0.3 (Jiang et al., 2023) (3-seed) confirms the PPL advantage (-17.3%) with directional commonsense gains but tied MMLU; preliminary evaluation on Qwen-2.5-14B (Team, 2024) shows consistent PPL improvement but reasoning decline (GSM8K -19.9pp , $\sigma=5.3\text{pp}$; MMLU -2.8pp). This paper is primarily an empirical study characterizing how the training objective shapes sparsity prediction quality across allocation strategies, sparsity levels, and model scales. The central finding—that output-distribution alignment consistently and substantially outperforms per-neuron classification for learned sparsity—is the paper’s contribution. Without sparse kernel support, the predictor adds 6.4% latency overhead; wall-clock speedup is not addressed (Appendix A). While individual components (KL divergence, Gumbel-Sigmoid relaxation) are established techniques, their application to contextual activation sparsity prediction has not been explored, and the resulting empirical behaviors—KL-discovered masks diverge 46–68% from magnitude targets yet reduce PPL by 14.5%—are not predictable from the component properties alone.

Our contributions are:

- *Empirical*: Output-distribution alignment ranking across five losses, with per-layer decomposition showing the signal type gap (9.47 PPL) exceeds the coupling breadth range (2.08 PPL), though confounded by train–inference mismatch (§4.5);
- *Empirical*: Cross-scale evidence (7–14B) with task-type split: KL outperforms on PPL and commonsense but shows systematic knowledge regression, structural to output-distribution optimization (C4 control; §4);
- *Diagnostic*: Target-swap ablation suggests KL-discovered mask targets transfer to BCE; per-layer controls suggest task-dependent coupling effects (§4.5);
- *Practical*: KL-derived masks render post-hoc compensation (+67.2M params) redundant at both tested sparsity levels (§4.5).

2 Related Work

Training-Free Activation Sparsity. TEAL (Liu et al., 2025) applies global top- k selection by activa-




Figure 1: Method overview: a lightweight per-layer predictor produces importance scores trained via KL divergence with Gumbel-Sigmoid relaxation. At inference, scores feed TEAL’s global top- k allocation without compensation networks.

tion magnitude and achieves 40–50% sparsity with no training cost. CATS (Lee et al., 2024) adapts per-layer thresholds for SiLU/SwiGLU activations; R-Sparse (Zhang et al., 2025b) predicts sparsity from rank-aware SVD projections; DIP (Federici et al., 2024) combines dynamic input pruning with cache-aware masking; SparseInfer (Shin et al., 2025b) and GRASP (Shin et al., 2025a) predict activation sparsity patterns without additional training. LaRoSA (Huang et al., 2025) applies layerwise orthogonal rotations for structured activation sparsity with wall-clock speedup on SiLU models. Wang et al. (Wang et al., 2025), AWQ (Lin et al., 2024), and Outlier Suppression+ (Wei et al., 2023) address sparsity–quantization interactions and activation-aware quantization. Dhar et al. (Dhar et al., 2024) characterize activation sparsity opportunities across architectures, showing that SwiGLU models exhibit lower natural sparsity than ReLU counterparts. OWL (Yin et al., 2024) determines per-layer sparsity ratios from outlier distributions for weight pruning. ActTail (Zheng et al., 2026) applies heavy-tailed self-regularization for global activation sparsity allocation; WiSparse (Yang, 2026) proposes weight-aware mixed activation sparsity. All training-free methods determine masks from fixed statistics or input-independent criteria, while our predictor generates *input-dependent* masks trained to minimize output-distribution divergence.

Predictor-Based Contextual Sparsity. DejaVu (Liu et al., 2023) introduced learned sparsity predictors trained via BCE to classify each neuron independently. ShadowLLM (Akshauri et al., 2024) uses attention-based routing with per-layer MSE supervision for intermediate coupling within each layer but without end-to-end gradient signal. FastForward (Gautam et al., 2026) adds a 67.2M-parameter compensation network using BCE for mask prediction, adding inference latency. SVD-CSP (Serbin et al., 2026) compresses sparsity

predictors via SVD decomposition for reduced predictor overhead. Dhar et al. (Dhar et al., 2025) compress activation patterns via clustering for efficient neuron-level prediction. LazyLLM (Fu et al., 2024) dynamically prunes tokens across layers for long-context efficiency; COTS (Zhang et al., 2025a) exploits connected activation patterns for structured sparse inference. PowerInfer (Song et al., 2024a) and PowerInfer-2 (Xue et al., 2024) exploit power-law activation frequencies for GPU-CPU/NPU scheduling. MaskLLM (Fang et al., 2024) shares the recipe-level principle of frozen weights + learnable masks + KL training. The methods differ in three respects: (a) MaskLLM’s mask parameters are fixed per input (static N:M weight sparsity via Gumbel-Softmax), while our predictor generates *per-token* masks that must generalize to unseen inputs; (b) static N:M patterns have no cross-layer allocation interaction, whereas our masks coordinate through TEAL’s global top- k ; (c) activation masks must be robust to distribution shift at inference, while weight masks are input-independent. Our contribution beyond the shared recipe is analytical: a systematic decomposition of signal type vs. coupling scope for sparsity prediction training, via target-swap ablation and per-layer KL controls (§4.5). All prior predictor-based methods for activation sparsity use decomposed (BCE) or intermediate (MSE) objectives where each neuron’s gradient is independent of other mask states; no prior work systematically compares loss functions for contextual sparsity prediction quality. We empirically observe that an output-distribution-coupled objective propagates gradients through all downstream layers and quantify when this matters for global allocation quality.

Static Bias Correction. SPON (Spontaneous Neurons; Xu et al., 2025) learns input-independent bias vectors via closed-form KL minimization per layer, absorbed into weights at deployment. Our method is fundamentally different: a per-token contextual predictor trained end-to-end across all layers that produces input-dependent masks. SPON solves a static per-layer problem; we solve a global dynamic prediction problem. The two are weakly correlated (cosine similarity 0.196, $R^2=0.038$; Appendix Q).

Activation Function Engineering. ReLU Strikes Back (Mirzadeh et al., 2024) replaces SiLU with ReLU for >90% natural sparsity; Tur-

boSparse (Song et al., 2024b) and ProSparse (Song et al., 2025) engineer extreme sparsity through activation modifications; the Sparsing Law (Luo et al., 2024) formalizes how activation density increases with training compute. MiniTron (Muralidharan et al., 2024) combines structured pruning with distillation but modifies architecture permanently. Haziza et al. (Haziza et al., 2025) demonstrate 2:4 semi-structured activation sparsity with hardware-native acceleration, complementing our unstructured approach. MoC (Wu et al., 2025b) selectively activates Top- K FFN channels using SwiGLU’s native gating during pretraining. All require retraining or weight modification; our method is a plug-in predictor on *unmodified* SwiGLU models.

Weight Pruning, Distillation, and Routing. Wanda (Sun et al., 2024), SparseGPT (Frantar and Alistarh, 2023), and SliceGPT (Ashkboos et al., 2024) produce static weight masks via one-shot or structured pruning; Polar Sparsity (Shrestha et al., 2025) scales contextual sparsity to batched inference by shifting focus from MLP to attention head sparsity; Alizadeh et al. (Alizadeh et al., 2024) exploit activation sparsity for memory-efficient inference via selective flash loading. Switch Transformers (Fedus et al., 2022) and Mixture-of-Depths (Rapooso et al., 2024) achieve conditional computation through learned routing. KL-based distillation (Kim and Rush, 2016; Sanh et al., 2019; Gu et al., 2024; Agarwal et al., 2024; Ko et al., 2024) trains compact students; SparseLoRA (Khaki et al., 2025) exploits contextual sparsity for fine-tuning. We apply the KL principle to train a sparsity predictor on unmodified dense models, using Gumbel-Sigmoid sampling (Maddison et al., 2017; Louizos et al., 2018) for differentiable discrete mask optimization.

MoE Routing and Gated Sparsity. Expert Choice routing (Zhou et al., 2022) lets each expert select its top- k tokens under a fixed capacity budget, a global allocation perspective that parallels TEAL’s neuron-level global top- k . D2D-MoE (Szatkowski et al., 2024) converts dense SwiGLU layers into dynamic- k MoE by exploiting the same activation sparsity we target, but restructures the network into discrete experts rather than predicting masks on the original architecture. MoNE (Cheng et al., 2025) applies neuron-granular top- k selection within each MoE expert during pretraining for fine-grained conditional computation without

additional routing parameters.

3 Method

We present an output-distribution KL framework for training contextual sparsity predictors in SwiGLU-based language models (Figure 1). The method replaces per-neuron binary classification (BCE) with an output-distribution alignment objective that couples all pruning decisions through the model’s forward pass and achieves competitive performance without auxiliary error compensation networks at inference.

3.1 Problem Formulation

Consider a pretrained transformer with L layers, each containing a SwiGLU feed-forward network (Shazeer, 2020). For layer ℓ , the FFN computes:

$$\mathbf{y}^{(\ell)} = (\sigma(\mathbf{x}^{(\ell)} \mathbf{W}_g^{(\ell)}) \odot \mathbf{x}^{(\ell)} \mathbf{W}_u^{(\ell)}) \mathbf{W}_d^{(\ell)}, \quad (1)$$

where $\mathbf{x}^{(\ell)} \in \mathbb{R}^d$ is the hidden state after attention ($d=4096$), σ denotes SiLU, $\mathbf{W}_g^{(\ell)}, \mathbf{W}_u^{(\ell)} \in \mathbb{R}^{d \times d_{\text{ff}}}$ and $\mathbf{W}_d^{(\ell)} \in \mathbb{R}^{d_{\text{ff}} \times d}$ are the gate, up-projection, and down-projection matrices, and $d_{\text{ff}}=14,336$.

Define the gate activation $\mathbf{g}^{(\ell)} = \sigma(\mathbf{x}^{(\ell)} \mathbf{W}_g^{(\ell)}) \in \mathbb{R}^{d_{\text{ff}}}$. Each element $g_j^{(\ell)}$ determines the contribution of neuron j to the output. We introduce a binary mask $\mathbf{m}^{(\ell)} \in \{0, 1\}^{d_{\text{ff}}}$ that zeros out pruned neurons:

$$\mathbf{y}_{\text{sparse}}^{(\ell)} = (\mathbf{m}^{(\ell)} \odot \mathbf{g}^{(\ell)} \odot \mathbf{x}^{(\ell)} \mathbf{W}_u^{(\ell)}) \mathbf{W}_d^{(\ell)}. \quad (2)$$

Our goal is to learn a mask predictor $f_{\theta}^{(\ell)}(\mathbf{x}^{(\ell)}) \rightarrow \mathbf{s}^{(\ell)} \in (0, 1)^{d_{\text{ff}}}$ that produces per-neuron importance scores from $\mathbf{x}^{(\ell)}$. At inference, the scores are fed to a global top- k allocation mechanism (Liu et al., 2025) that distributes the sparsity budget across layers non-uniformly, retaining the $(1-\rho) \times L \times d_{\text{ff}}$ highest-scoring neurons model-wide for a target sparsity ratio ρ .

3.2 Predictor Architecture

Each layer ℓ has an independent two-layer MLP predictor:

$$f_{\theta}^{(\ell)}(\mathbf{x}^{(\ell)}) = \text{GELU}(\mathbf{x}^{(\ell)} \mathbf{W}_1^{(\ell)}) \mathbf{W}_2^{(\ell)}, \quad (3)$$

where $\mathbf{W}_1^{(\ell)} \in \mathbb{R}^{d \times d_b}$ and $\mathbf{W}_2^{(\ell)} \in \mathbb{R}^{d_b \times d_{\text{ff}}}$ with bottleneck dimension $d_b=128$. The predictor takes $\mathbf{x}^{(\ell)}$ (the hidden state after attention, before the

FFN) and outputs d_{ff} -dimensional logits $\mathbf{s}^{(\ell)}$, from which the mask is derived.

Each predictor has $(4096 \times 128) + (128 \times 14336) = 2.37\text{M}$ parameters. With 32 independent predictors (one per layer), the total overhead is 75.9M parameters, less than 1% of the 8B base model. The predictors operate independently across layers, requiring no cross-layer communication at inference.

3.3 Gumbel-Sigmoid Relaxation

The hard mask $m_j = \mathbf{1}[s_j > \text{threshold}]$ is non-differentiable, preventing gradient flow from the KL objective through the mask to the predictor parameters. We apply Gumbel-Sigmoid relaxation (Maddison et al., 2017; Louizos et al., 2018; Jang et al., 2017) to produce differentiable soft masks:

$$\tilde{m}_j^{(\ell)} = \sigma_{\tau} \left(\log \frac{s_j^{(\ell)}}{1-s_j^{(\ell)}} + \varepsilon_1 - \varepsilon_2 \right), \quad \varepsilon_{1,2} \sim \text{Gumbel}(0, 1), \quad (4)$$

where $\sigma_{\tau}(x) = \text{sigmoid}(x/\tau)$ and τ is a temperature parameter. The Gumbel noise provides stochastic exploration of mask configurations during training, while the temperature controls sharpness: as $\tau \rightarrow 0$, $\tilde{m}_j$ approaches a binary decision.

We anneal τ linearly from 1.0 to 0.1 over training. At high temperature, the soft mask permits gradient flow through all neurons for broad exploration. As temperature decreases, the mask concentrates on confident decisions, approximating the hard allocation used at inference.

Comparison with Straight-Through Estimation.

The straight-through estimator (STE) (Ben-gio et al., 2013) applies a hard threshold in the forward pass and passes gradients through unchanged ($\partial m / \partial s = 1$). In our setting, STE produces discontinuous gradient signals: small parameter changes can flip many neurons simultaneously. The resulting oscillations in the penalty weight prevent convergence to the target sparsity. Gumbel relaxation eliminates this pathology and contributes a 33.8% improvement (PPL 15.77 $\rightarrow$ 11.79 at 50% sparsity), as we verify empirically in §4.

3.4 KL Divergence Training Objective

Let $\mathbf{z}_{\text{dense}}$ denote the output logits from a full dense forward pass, and $\mathbf{z}_{\text{sparse}}$ the logits from a forward pass with predictor-generated soft masks $\{\tilde{\mathbf{m}}^{(\ell)}\}_{\ell}$ applied at every layer. We minimize the forward KL divergence between the dense and sparse output

distributions:

$$\mathcal{L}_{\text{KL}} = D_{\text{KL}}(p_{\text{dense}} \| p_{\text{sparse}}) = \sum_v p_{\text{dense}}^{(v)} \log \frac{p_{\text{dense}}^{(v)}}{p_{\text{sparse}}^{(v)}}, \quad (5)$$

where $p_{\text{dense}} = \text{softmax}(\mathbf{z}_{\text{dense}})$ and $p_{\text{sparse}} = \text{softmax}(\mathbf{z}_{\text{sparse}})$, summed over the vocabulary v and averaged over all token positions in the sequence. The base model is frozen; only the predictor parameters receive gradients.

We use forward KL (mass-covering) rather than reverse KL, as the goal is to preserve the dense model’s primary predictions rather than explore alternative modes. In sparsity prediction, mode-covering is critical: false negatives (omitting a rarely but critically active neuron) cascade through downstream layers under global allocation, whereas false positives merely waste capacity; reverse KL’s mode-seeking concentrates on high-probability neurons and risks missing compositionally important activations (Wu et al., 2025a). §4 validates this choice empirically: forward KL outperforms reverse KL by 8.6% and JSD by 6.4% (at 50% sparsity; the gap narrows at 30%; Table 4). No temperature scaling ($T=1$) is applied (Hinton et al., 2015); we distinguish T from the Gumbel temperature τ (§3.3), which controls mask sharpness rather than logit scaling.

3.5 Geometric Penalty Schedule

We enforce a model-level sparsity target ρ via constrained optimization with $\bar{s} = \frac{1}{L} \sum_{\ell} \frac{1}{d_{\text{ff}}} \sum_j \tilde{m}_j^{(\ell)}$ denoting the average retained fraction:

$$\min_{\theta} \mathcal{L}_{\text{KL}} \quad \text{s.t.} \quad \bar{s} = 1 - \rho. \quad (6)$$

We relax this via a geometric penalty schedule:

$$\mathcal{L} = \mathcal{L}_{\text{KL}} + \lambda \cdot \text{clamp}(|\bar{s} - (1 - \rho)| - \epsilon, 0)^2, \quad (7)$$

where $\epsilon = 0.02$ is a deadzone margin that suppresses oscillation near the target (sensitivity in Appendix C). The penalty weight λ is updated via geometric doubling (doubled when sparsity exceeds target, halved otherwise), clamped to $[\lambda_{\text{floor}}, \lambda_{\text{max}}]$, reaching target sparsity within 40–110 steps across all tested regimes (Appendix C).

3.6 Gradient Analysis: BCE vs. KL

We contrast the gradient structure of BCE and KL to build intuition for how KL discovers divergent mask patterns (§4.4). This analysis is suggestive;

Table 1: Taxonomy of learned sparsity prediction methods along correction granularity (rows) and training signal coupling (columns). Our method occupies the Per-token $\times$ Coupled cell (organizational framework). Empty cells represent unexplored combinations; their relative merit is an open question.

	Decomposed (BCE)	Intermediate (MSE)	Coupled (KL)
Static	—	—	SPON, MaskLLM
Per-token	DejaVu, FastForward	ShadowLLM, D2DMoE	Ours

empirical evidence in §4.4 shows KL-discovered mask targets drive inference quality.

Under BCE, the predictor trains against oracle masks $m_j^* = \mathbf{1}[|g_j| \geq \text{top-}k \text{ threshold}]$ derived from activation magnitudes. The gradient $\partial \mathcal{L}_{\text{BCE}} / \partial s_j = \sigma(s_j) - m_j^*$ depends only on neuron j ; each neuron is trained independently with no inter-neuron coupling. Under output-distribution KL, the chain rule introduces a downstream Jacobian $\partial \mathbf{z}_{\text{sparse}} / \partial \tilde{m}_j^{(\ell)}$ that couples neuron j ’s gradient to all other retained neurons in layers ℓ through L , because the sparse hidden state propagates through subsequent attention and FFN operations. This coupling allows KL to evaluate each neuron’s importance in the context of the full sparse computation, rather than against a fixed magnitude target.

Gradient SNR measurements¹ reveal that BCE maintains consistently higher SNR than KL throughout training ($2.3\times$ at convergence; Figure 6; extended analysis in Appendix H), indicating that KL’s quality advantage does not arise from cleaner gradient signals. Rather, KL’s end-to-end coupled gradient serves as a *search mechanism* that discovers compositionally superior mask patterns, even when the per-step signal is noisier than BCE’s decomposed updates. The target-swap ablation in §4.4 suggests that the discovered mask targets, not the coupling mechanism itself, drive inference-stage quality.

3.7 Two-Dimensional Taxonomy of Sparsity Prediction Methods

We organize prior work along two axes: *correction granularity* (how the mask adapts to input) and *training signal coupling* (whether the objective captures inter-neuron dependencies).

Our method combines per-token prediction with coupled KL training (Table 1). KL on output logits directly optimizes next-token prediction loss, which composes across layers; per-layer MSE does not—a mask minimizing FFN output error at layer ℓ may amplify errors in layers $\ell+1 \dots L$,

¹ $\text{SNR}_{\ell} = \|\mathbb{E}[\nabla_{\theta_{\ell}} \mathcal{L}]\|_2^2 / \text{Var}[\nabla_{\theta_{\ell}} \mathcal{L}]$, expectation over mini-batches.

explaining its $1.7\times/9.6\times$ worse PPL than TEAL at 30%/50% (50%: s42 only; §4). Full taxonomy discussion in Appendix O.

4 Experiments

4.1 Experimental Setup

Model and Predictor. We evaluate on LLaMA-3.1-8B (Grattafiori et al., 2024), which uses SwiGLU FFN blocks (Shazeer, 2020) with $d_{\text{model}} = 4096$ and $d_{\text{ff}} = 14336$ across 32 layers. The sparsity predictor is a two-layer MLP with 2.37M parameters per layer (75.9M total), implemented as $\text{Linear}(4096 \rightarrow 128) + \text{GELU} + \text{Linear}(128 \rightarrow 14336)$, trained to produce per-neuron importance scores that minimize output-distribution divergence under sparsity.

Training. All predictors are trained for 1000 steps on WikiText-103 with sequence length 2048 and batch size 4 on a single NVIDIA RTX PRO 6000 Blackwell Server Edition GPU (96 GB), using AdamW (Loshchilov and Hutter, 2019) with learning rate 1×10^{-3} , weight decay 1×10^{-2} , and cosine decay. The KL predictor uses Gumbel-Sigmoid relaxation (Jang et al., 2017) with temperature $\tau = 1.0$ annealed to 0.1 over training.

Evaluation. We report perplexity on WikiText-2 (Merity et al., 2017) (seq_len=2048) and downstream accuracy on ARC-Challenge (25-shot), WinoGrande (5-shot), MMLU (5-shot), GSM8K (8-shot), HellaSwag (10-shot), and PIQA (0-shot).

Inference Protocol. At inference, the predictor produces per-neuron scores that are passed to TEAL’s global top- k allocation mechanism (Liu et al., 2025), which distributes the sparsity budget across layers non-uniformly. We term this *predictor-TEAL*: the predictor replaces TEAL’s magnitude-based scoring while retaining its allocation strategy.

Baselines. We compare against: (1) **Vanilla TEAL** (Liu et al., 2025), which uses activation magnitudes directly²; (2) **BCE predictor** (DejaVu-style (Liu et al., 2023)), trained with per-neuron binary cross-entropy on magnitude-based oracle masks for 1000 and 2700 steps; and (3) **Uniform allocation**, which distributes the sparsity budget equally across layers.

²TEAL is deterministic; single-run evaluation suffices. KL/BCE predictors require multi-seed evaluation due to random initialization.

4.2 Main Results

Table 2 presents the primary comparison. Uniform allocation preserves downstream performance best; global allocation is analyzed to understand allocation-quality tradeoffs (§4.6). At 30% sparsity, output-distribution KL under uniform allocation achieves PPL 8.81 ± 0.02 (3-seed), a 14.5% reduction over TEAL at matched actual sparsity (8.81 vs. 10.30, both at 28.5%; Appendix D). Uniform allocation yields stable downstream (MMLU $40.7 \pm 1.3\%$, 3-seed; Appendix F), while global top- k exhibits high seed variance ($35.7 \pm 9.5\%$) from late-layer starvation (L31 retention 28.4%, Appendix E; §4.6).

Task-Type Split. KL shows directional improvements over TEAL on commonsense reasoning (ARC-c, WG, HellaSwag, PIQA: 4/4 wins at 30% uniform; Appendix I) but underperforms on knowledge-intensive tasks (MMLU, GSM8K: 0/2). With only three seeds, 95% confidence intervals for individual downstream tasks overlap; we therefore treat PPL ($\sigma=0.02$, non-overlapping) as the primary metric and downstream differences as directional evidence. The KL predictor protects L14–L18 (14–18% sparsity vs. 32% global mean; Appendix U) while aggressively pruning L22–L28 (48–57%). Constraint interventions (Appendix T) confirm a task-dependent trade-off: removing the valley improves MMLU +3.4pp but reduces ARC-c by 1.5pp, consistent with localized factual associations (Geva et al., 2023); task-aware allocation may resolve this.

Knowledge-Task Regression: Hypothesis. Constraint experiments (Appendix T) show allocation contributes: removing the learned valley improves MMLU +3.4pp and GSM8K +1.6pp. A C4-trained control (Appendix V) does not alleviate the regression (MMLU 40.0%, GSM8K 4.1% vs. 40.7%, 7.5%), suggesting the deficit is structural to the optimization objective rather than training data composition. Both loss-inherent (mode-covering sacrifices rare activations) and allocation-induced factors likely contribute; disentangling them requires neuron-level attribution beyond our scope.

At 50%, the PPL advantage widens (11.92 vs. 23.90, 50% reduction) and generalizes to C4 (Appendix D), but all methods collapse to near-random on MMLU ($\leq 29\%$) and GSM8K ($\leq 2.7\%$); no allocation strategy recovers at this capacity limit.

Table 2: Main results on LLaMA-3.1-8B. * $2.71\times$ KL compute (§4.3). $\pm$: 3-seed; unmarked: s42. † Actual sparsity: KL 28.5% vs. TEAL 29.3% at 30% (Appendix D).

Sparsity	Method	PPL↓	ARC-c	WG	MMLU	GSM8K
0%	Dense	6.12	58.28	78.06	65.27	48.75
30%	TEAL uniform	10.82	39.33	65.04	45.86	8.79
	BCE global	11.25	41.2	66.0	36.3	6.0
	BCE uniform	15.09 $\pm$ 0.13	33.45	61.09	43.05	3.97
	MSE uniform	18.52 $\pm$ 1.68	35.8 $\pm$ 4.1	59.8 $\pm$ 2.6	30.7 $\pm$ 2.6	2.3 $\pm$ 0.7
	KL global (Ours) †	8.96 $\pm$ 0.02	43.8$\pm$3.6	70.2$\pm$0.9	35.7 $\pm$ 9.5	8.0 $\pm$ 0.7
	KL uniform †	8.81$\pm$0.02	43.7 $\pm$ 2.5	68.9 $\pm$ 0.4	40.7 $\pm$ 1.3	7.5 $\pm$ 0.1
50%	TEAL uniform	23.90	27.47	54.38	29.26	1.74
	BCE (2700 steps)*	28.13 $\pm$ 0.20	27.6 $\pm$ 0.6	51.8 $\pm$ 0.6	26.1 $\pm$ 0.5	2.1 $\pm$ 0.2
	MSE uniform	229.26	25.00	50.59	23.79	1.59
	KL global (Ours)	11.84$\pm$0.15	36.3$\pm$3.2	63.8$\pm$0.3	28.3 $\pm$ 0.7	2.7$\pm$1.3
	KL uniform	11.92 $\pm$ 0.07	33.8 $\pm$ 1.2	62.9 $\pm$ 0.1	28.0 $\pm$ 0.6	2.2 $\pm$ 0.1

4.3 Extended Training Analysis

At $2.71\times$ KL compute budget (2700 steps, matched wall-clock), BCE achieves PPL 28.13 $\pm$ 0.20 (3-seed), still $2.39\times$ worse than KL (11.79). The gap is systematic: BCE reaches higher oracle F1 (0.613 vs. 0.540) with 139% worse PPL, confirming KL optimizes for distributional fidelity rather than magnitude agreement.

4.4 Mask Pattern Analysis

KL-derived masks diverge from magnitude oracles: per-layer Jaccard distance averages 68% at 50% (64 samples, 131K tokens; Figure 8), with cross-seed variation (36.3%) confirming systematic divergence (Appendices F.1, M). Despite this, KL achieves PPL 11.79 vs. BCE’s 28.13 $\pm$ 0.20, confirming that distributional fidelity, not magnitude agreement, produces better masks. Divergence follows a U-shaped per-layer pattern, peaking in late layers (0.70–0.75 at 50%).

Target-Swap Ablation: Target Quality vs. Gradient Coupling. To disentangle whether KL’s advantage stems from its coupled gradient pathway or from the mask patterns it discovers, we train a standard BCE predictor on KL-derived oracle masks instead of magnitude-based masks (C1 target-swap ablation). Table 3 reports results under TEAL uniform allocation at matched sparsity.

At 50% (s42, earlier oracle), C1 matches KL within 0.79pp, while magnitude-trained BCE yields $2.4\times$ worse PPL (Table 2). At 30%, retrained seeds show ≤ 1.18 pp agreement³. C1 demonstrates that KL-discovered mask patterns *transfer* to a simpler training mechanism (PPL within 0.79pp; downstream variation up to 4.5pp MMLU); per-layer KL controls (§4.5) further suggest task-dependent coupling effects. However, C1 also changes mask

³s42 oracle from earlier checkpoint (gap 0.65pp); s123/s456 use seed-matched converged oracles (≤ 0.03 pp gap).

Table 3: Target-swap ablation (C1): BCE on KL-derived masks vs. KL predictor (same seed), TEAL uniform. All seeds show max $\Delta \leq 1.70$ pp at 30%; max $\Delta \leq 0.79$ pp at 50% (s42). † **Provenance:** s42 oracle extracted from an earlier training run (not seed-matched); s123/s456 use seed-matched retrained oracles (Appendix P).

Sparsity	Seed	Method	PPL↓	ARC-c	WG	MMLU	GSM8K
30%	s42 †	C1	9.45	40.27	67.40	42.30	7.58
		KL	8.82	40.78	68.59	40.6	7.35
	s123	C1	8.84	44.71	69.77	39.92	6.82
		KL	8.83	45.31	69.38	39.4	7.51
	s456	C1	8.82	44.71	69.93	41.04	6.60
		KL	8.79	45.05	68.75	42.0	7.58
50%	s42	C1	11.96	32.17	63.14	27.79	2.05
		KL	11.96	32.59	62.35	27.92	1.97

Table 4: Ablation on LLaMA-3.1-8B (WikiText-2 PPL↓; uniform at 30%, global at 50%). $\pm$: 3-seed; unmarked: s42. † STE: 20.8% actual sparsity at 30% target; loss effect (Δ PPL=20.4) is $5.1\times$ relaxation effect (Δ PPL=3.98). § MSE 50%: 2/3 seeds collapsed to $\sim 5\%$ actual sparsity; s42 only (Appendix G).

Configuration	30% PPL↓	50% PPL↓	Δ_{50}
<i>KL-based variants</i>			
Forward KL + Gumbel (Ours)	8.81$\pm$0.02	11.79	—
JSD (symmetric) + Gumbel	8.90 $\pm$ 0.02	12.54	+0.75
Reverse KL + Gumbel	9.26 $\pm$ 0.07	12.80	+1.01
Forward KL + STE	— †	15.77	+3.98
<i>Per-layer training</i>			
Per-layer MSE	18.52 $\pm$ 1.68	229.26 §	+217.47
Per-layer KL ($K=1$)	9.05 $\pm$ 0.02	34.85	+23.06
Per-layer KL ($K=4$)	8.89 $\pm$ 0.01	44.23	+32.44
<i>BCE-based variants</i>			
BCE (no compensation)	15.71	27.97	+16.18
BCE + Gumbel (no penalty)	—	32.22	+20.43
BCE + Gumbel + Penalty	—	38.02	+26.23
BCE + STE	—	60.54	+48.75

representation: binarizing oracle soft logits (>0) into hard masks discards ranking information, so the near-zero delta reflects transferability of binary mask *patterns* rather than full predictor knowledge. BCE+Gumbel (32.22) underperforms standard BCE (27.97), indicating trajectory-dependent confounds that prevent clean isolation of target quality from gradient coupling (Appendix P).

4.5 Ablation Studies

Table 4 isolates the contribution of each design choice at 30% and 50% sparsity.

Relaxation, Divergence, and SPON. At 50%, Gumbel-Sigmoid outperforms STE (Bengio et al., 2013) by 33.8% (11.79 vs. 15.77); forward KL outperforms JSD (+6.4%) and reverse KL (+8.6%), consistent with mode-covering (Gu et al., 2024). At 30%, forward KL (8.81) and JSD (8.90) are near-identical, both outperforming reverse KL (9.26) and MSE (18.52 $\pm$ 1.68). KL–JSD near-equivalence suggests the divergence measure matters less than optimizing through the output distribution; cou-

pling scope may provide additional downstream benefits (§4.5).

Per-Layer KL: Empirical Signal Type vs. Coupling Scope. To isolate cross-layer coupling, we train per-layer KL predictors that mask K random layers per step and compute KL on the full model output, preserving the output-distribution signal without cross-layer gradients. The K -curve (Figure 2) reveals three empirical patterns at 30%: (1) PPL is non-monotonic, with $K=2$ worst (10.75 ± 0.15) due to insufficient per-layer updates and $K=16$ best (8.67 ± 0.02), marginally outperforming E2E (8.81 ± 0.02); (2) at $K=1$ (no coupling), per-layer KL achieves 9.05 ± 0.02 , within 2.7% of E2E, while per-layer MSE degrades to 18.52 ± 1.68 ($2.05 \times$ worse); (3) the signal type⁴ gap ($K=1$ KL 9.05 vs. MSE 18.52: 9.47 PPL) exceeds the coupling breadth range ($K=16$ 8.67 vs. $K=2$ 10.75: 2.08 PPL). This suggests output-distribution alignment contributes more than coupling scope to PPL, but the decomposition is confounded by train–inference mismatch ($K=1$ masks 1/32 layers during training vs. all 32 at inference), per-layer update count disparity, and the 50% failure where $K=1$ (34.85) outperforms $K=4$ (44.23) while E2E reaches 11.84. A 3-seed $K=4$ control⁵ partially deconfounds (Appendix S): PPL and GSM8K gaps largely close while commonsense gaps (ARC-c, WG) persist, suggesting task-dependent coupling effects; detailed per-task attribution is in Appendices S–R; full K-curve downstream in Appendix W.

Compensation Redundancy. FastForward-style compensation (Gautam et al., 2026) (+67.2M params) provides no benefit (Appendices J–B.2); SPON bias correction is orthogonal (Appendix Q). Predictor capacity and temperature sensitivity in Appendix B.

4.6 Sparsity Scaling and Regime Analysis

Allocation Behavior. Table 14 (Appendix L) reveals non-monotonic allocation preference: uniform is best at $\leq 30\%$ ($7.3 \times$ smaller MMLU σ), global at 35–40%, both collapse at $\geq 45\%$.

⁴We use “signal type” as shorthand for the training objective’s optimization target (output distribution vs. intermediate activations); this dimension also entails differences in gradient flow topology and optimization landscape that are not independently controlled.

⁵Actual sparsity 29.6%; per-layer sparsity std=0.003 (near-uniform).




Figure 2: Coupling breadth K vs. perplexity (30%, TEAL uniform, 3-seed mean $\pm$ std). $K=32$ denotes end-to-end training. $K=2$ degradation reflects insufficient per-layer updates (~ 62 vs. ~ 250 at $K=8$).




Figure 3: PPL scaling with sparsity (global allocation). KL’s advantage over TEAL grows with sparsity.

Cross-Architecture and Scale. On Mistral-7B-v0.3 (Jiang et al., 2023), PPL generalizes (-17.3% , 3-seed; Appendix N); on Qwen-2.5-14B (Team, 2024), PPL generalizes but MMLU/GSM8K decline (Appendix K). Consistent regression across scales (8B MMLU -5.2 pp, 14B GSM8K 32.2 ± 5.3 vs. 52.1) suggests systematic limitations of output-distribution KL. The pipeline adds 6.4% latency overhead absent sparse kernels (Appendix A).

5 Conclusion

Output-distribution alignment is the primary design dimension for learned sparsity prediction, outperforming BCE and MSE across five losses. The KL predictor reduces PPL by 14.5% over TEAL at 30%, confirmed on Mistral-7B (-17.3%) and Qwen-2.5-14B. Gains are task-type dependent: KL improves commonsense reasoning but degrades knowledge-intensive tasks, suggesting KL for commonsense/PPL and TEAL for knowledge tasks (§4). Downstream differences are directional ($n=3$); code released upon acceptance.

Limitations

Our findings are bounded by four limitations. **Spar-sity sweet spot.** 30% is the lowest-degradation level; at 50%, all methods collapse on reasoning tasks, and a separate predictor must be retrained per target. **Knowledge-intensive regression.** KL consistently underperforms on MMLU and GSM8K across scales (8B: -5.2pp MMLU; 14B: -19.9pp GSM8K), suggesting output-distribution objectives systematically sacrifice reasoning-critical path-ways. **Evaluation scope.** Open-ended generation (e.g., MT-Bench) and contexts beyond 2048 tokens remain untested. **Baseline coverage.** The 30% isocompute comparison lacks a BCE 2700-step baseline; WikiText-2 evaluation shares provenance with the training corpus (WikiText-103), partially mitigated by C4 OOD results (Appendix D). MSE uses identical hyperparameters to KL; its gradient magnitude is $\sim 5\times$ larger, and performance may improve with a dedicated learning rate.

References

- Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos, Matthieu Geist, and Olivier Bachem. 2024. [On-policy distillation of language models: Learning from self-generated mistakes](#). In *Proceedings of the International Conference on Learning Representations (ICLR)*.
- Yash Akhauri, Ahmed F AbouElhamayed, Jordan Dotzel, Zhiru Zhang, Alexander M Rush, Safeen Huda, and Mohamed S Abdelfattah. 2024. [Shad-owlm: Predictor-based contextual sparsity for large language models](#). In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*.
- Keivan Alizadeh, Iman Mirzadeh, Dmitry Belenko, Karen Khatamifard, Minsik Cho, Carlo C Del Mundo, Mohammad Rastegari, and Mehrdad Farajtabar. 2024. LLM in a flash: Efficient large language model inference with limited memory. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*.
- Saleh Ashkboos, Maximilian L. Croci, Marcelo Gennari do Nascimento, Torsten Hoefler, and James Hensman. 2024. [Slicegpt: Compress large language models by deleting rows and columns](#). In *Proceedings of the International Conference on Learning Representations (ICLR)*.
- Yoshua Bengio, Nicholas Léonard, and Aaron Courville. 2013. [Estimating or propagating gradients through stochastic neurons for conditional computation](#). *arXiv preprint arXiv:1308.3432*.
- Runxi Cheng, Yuchen Guan, Yucheng Ding, Qingguo Hu, Yongxian Wei, Chun Yuan, Yelong Shen, Weizhu Chen, and Yeyun Gong. 2025. [Mixture of neuron experts](#). *arXiv preprint arXiv:2510.05781*.
- Nobel Dhar, Bobin Deng, Md Romyull Islam, Kazi Fahim Ahmad Nasif, Liang Zhao, and Kun Suo. 2024. [Activation sparsity opportunities for compressing general large language models](#). *arXiv preprint arXiv:2412.12178*.
- Nobel Dhar, Bobin Deng, Md Romyull Islam, Xinyue Zhang, Kazi Fahim Ahmad Nasif, and Kun Suo. 2025. [A sparsity predicting approach for large language models via activation pattern clustering](#). In *Proceedings of the International Conference on Parallel Processing (Euro-Par)*.
- Gongfan Fang, Hongxu Yin, Saurav Muralidharan, Greg Heinrich, Jeff Pool, Jan Kautz, Pavlo Molchanov, and Xinchao Wang. 2024. [Maskllm: Learnable semi-structured sparsity for large language models](#). In *Advances in Neural Information Processing Systems (NeurIPS)*.
- Marco Federici, Davide Belli, Mart van Baalen, Amir Jalalirad, Andrii Skliar, Bence Major, Markus Nagel, and Paul Whatmough. 2024. [Efficient llm inference using dynamic input pruning and cache-aware masking](#). In *Proceedings of Machine Learning and Systems (MLSys)*.
- William Fedus, Barret Zoph, and Noam Shazeer. 2022. [Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity](#). *Journal of Machine Learning Research*, 23(120):1–40.
- Elias Frantar and Dan Alistarh. 2023. [Sparsegpt: Massive language models can be accurately pruned in one-shot](#). In *Proceedings of the International Conference on Machine Learning (ICML)*.
- Qichen Fu, Minsik Cho, Thomas Merth, Sachin Mehta, Mohammad Rastegari, and Mahyar Najibi. 2024. [LazyLLM: Dynamic token pruning for efficient long context LLM inference](#). *arXiv preprint arXiv:2407.14057*.
- Aayush Gautam, Mukul Gagrani, Junyoung Park, Mingu Lee, Chiris Lott, and Narasimha Reddy. 2026. [Fast forward: Accelerating llm prefill with predictive fnn sparsity](#). *arXiv preprint arXiv:2602.00397*.
- Mor Geva, Jasmijn Bastings, Katja Filippova, and Amir Globerson. 2023. [Dissecting recall of factual associations in auto-regressive language models](#). In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 12216–12235.
- Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aieleen Lakber, Aimin Chen, Ait Mokhtar Salma, and Ajay Menon. 2024. [The llama 3 herd of models](#). *arXiv preprint arXiv:2407.21783*.

771	Yuxian Gu, Li Dong, Furu Wei, and Minlie Huang. 2024.	Zichang Liu, Jue Wang, Tri Dao, Tianyi Zhou, Binhang	827
772	Minillm: Knowledge distillation of large language	Yuan, Zhao Song, Anshumali Shrivastava, Ce Zhang,	828
773	models . In <i>Proceedings of the International Confer-</i>	Yundong Tian, Christopher Re, and Beidi Chen.	829
774	<i>ence on Learning Representations (ICLR)</i> .	2023. Deja vu: Contextual sparsity for efficient llms	830
775	Daniel Haziza, Timothy Chou, Dhruv Choudhary, Luca	at inference time . In <i>Proceedings of the International</i>	831
776	Wehrstedt, Francisco Massa, Jiecao Yu, Geonhwa	<i>Conference on Machine Learning (ICML)</i> .	832
777	Jeong, Supriya Rao, Patrick Labatut, and Jesse Cai.	Ilya Loshchilov and Frank Hutter. 2019. Decoupled	833
778	2025. Accelerating transformer inference and training	weight decay regularization . In <i>Proceedings of the</i>	834
779	with 2:4 activation sparsity . <i>arXiv preprint</i>	<i>International Conference on Learning Representa-</i>	835
780	<i>arXiv:2503.16672</i> .	<i>tions (ICLR)</i> .	836
781	Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. 2015.	Christos Louizos, Max Welling, and Diederik P. Kingma.	837
782	Distilling the knowledge in a neural network . In	2018. Learning sparse neural networks through l_0	838
783	<i>NeurIPS 2014 Deep Learning Workshop</i> .	regularization . In <i>Proceedings of the International</i>	839
784	Fangcheng Huang and 1 others. 2025. LaRoSA: Ac-	<i>Conference on Learning Representations (ICLR)</i> .	840
785	celerating large language model inference with lay-	Yuqi Luo, Chenyang Song, Xu Han, Yingfa Chen, Chao-	841
786	erwise rotated sparse activation. <i>arXiv preprint</i>	jun Xiao, Xiaojun Meng, Liqun Deng, Jiansheng Wei,	842
787	<i>arXiv:2507.01299</i> .	Zhiyuan Liu, and Maosong Sun. 2024. Sparsing law:	843
788	Eric Jang, Shixiang Gu, and Ben Poole. 2017. Categori-	Towards large language models with greater activa-	844
789	cal reparameterization with gumbel-softmax . In <i>Pro-</i>	tion sparsity . <i>arXiv preprint arXiv:2411.02335</i> .	845
790	<i>ceedings of the International Conference on Learning</i>	Chris J. Maddison, Andriy Mnih, and Yee Whye Teh.	846
791	<i>Representations (ICLR)</i> .	2017. The concrete distribution: A continuous relax-	847
792	Albert Q. Jiang, Alexandre Sablayrolles, Arthur Men-	ation of discrete random variables . In <i>Proceedings of</i>	848
793	sch, Chris Bamford, Devendra Singh Chaplot, Diego	<i>the International Conference on Learning Representa-</i>	849
794	de las Casas, Florian Bressand, Gianna Lengyel, Guil-	<i>tions (ICLR)</i> .	850
795	laume Lample, and Lucile Saulnier. 2023. Mistral	Stephen Merity, Caiming Xiong, James Bradbury, and	851
796	7b . <i>arXiv preprint arXiv:2310.06825</i> .	Richard Socher. 2017. Pointer sentinel mixture mod-	852
797	Samir Khaki, Xiuyu Li, Junxian Guo, Ligeng Zhu,	els . In <i>Proceedings of the International Conference</i>	853
798	Konstantinos N. Plataniotis, Amir Yazdanbakhsh,	<i>on Learning Representations (ICLR)</i> .	854
799	Kurt Keutzer, Song Han, and Zhijian Liu. 2025.	Iman Mirzadeh, Keivan Alizadeh, Sachin Mehta, Carlo	855
800	Sparselora: Accelerating llm fine-tuning with con-	C Del Mundo, Oncel Tuzel, Golnoosh Samei, Mo-	856
801	textual sparsity . In <i>Proceedings of the International</i>	hammad Rastegari, and Mehrdad Farajtabar. 2024.	857
802	<i>Conference on Machine Learning (ICML)</i> .	Relu strikes back: Exploiting activation sparsity in	858
803	Yoon Kim and Alexander M. Rush. 2016. Sequence-	large language models . In <i>Proceedings of the Inter-</i>	859
804	level knowledge distillation . In <i>Proceedings of the</i>	<i>national Conference on Learning Representations</i>	860
805	<i>Conference on Empirical Methods in Natural Lan-</i>	<i>(ICLR)</i> .	861
806	<i>guage Processing (EMNLP)</i> .	Saurav Muralidharan, Sharath Turuvekere Sreenivas,	862
807	Jongwoo Ko, Sungnyun Kim, Tianyi Chen, and Se-	Raviraj Joshi, Marcin Chochowski, Mostofa Patwary,	863
808	Young Yun. 2024. Distillm: Towards streamlined dis-	Mohammad Shoneybi, Bryan Catanzaro, Jan Kautz,	864
809	tillation for large language models . In <i>Proceedings of</i>	and Pavlo Molchanov. 2024. Compact language mod-	865
810	<i>the International Conference on Machine Learning</i>	els via pruning and knowledge distillation . In <i>Ad-</i>	866
811	<i>(ICML)</i> .	<i>vances in Neural Information Processing Systems</i>	867
812	Donghyun Lee, Je-Yong Lee, Genghan Zhang,	<i>(NeurIPS)</i> .	868
813	Mo Tiwari, and Azalia Mirhoseini. 2024. Cats:	Colin Raffel, Noam Shazeer, Adam Roberts, Katherine	869
814	Contextually-aware thresholding for sparsity in large	Lee, Sharan Narang, Michael Matena, Yanqi	870
815	language models . <i>arXiv preprint arXiv:2404.08763</i> .	Zhou, Wei Li, and Peter J. Liu. 2020. Exploring the	871
816	Ji Lin, Jiaming Tang, Haotian Tang, Shang Yang, Wei-	limits of transfer learning with a unified text-to-text	872
817	Ming Chen, Wei-Chen Wang, Guangxuan Xiao,	transformer . <i>Journal of Machine Learning Research</i> ,	873
818	Xingyu Dang, Chuang Gan, and Song Han. 2024.	21(140):1–67.	874
819	AWQ: Activation-aware weight quantization for	David Raposo, Sam Ritter, Blake Richards, Timothy	875
820	LLM compression and acceleration . In <i>Proceedings</i>	Lillicrap, Peter Conway Humphreys, and Adam San-	876
821	<i>of Machine Learning and Systems (MLSys)</i> .	toro. 2024. Mixture-of-depths: Dynamically allocat-	877
822	James Liu, Pragaash Ponnusamy, Tianle Cai, Han Guo,	ing compute in transformer-based language models .	878
823	Yoon Kim, and Ben Athiwaratkun. 2025. Training-	<i>arXiv preprint arXiv:2404.02258</i> .	879
824	free activation sparsity in large language models .	Victor Sanh, Lysandre Debut, Julien Chaumond, and	880
825	In <i>Proceedings of the International Conference on</i>	Thomas Wolf. 2019. Distilbert, a distilled version	881
826	<i>Learning Representations (ICLR)</i> .	of bert: smaller, faster, cheaper and lighter . <i>arXiv</i>	882
		<i>preprint arXiv:1910.01108</i> .	883

884	Georgii Serbin, Kirill Koshkin, Zhongao Sun, Anastasiya Bistrigova, and C.C. Korikov. 2026. Svd contextual sparsity predictors for fast llm inference . <i>arXiv preprint arXiv:2603.14110</i> .	937
885		938
886		939
887		940
888	Noam Shazeer. 2020. Glu variants improve transformer . <i>arXiv preprint arXiv:2002.05202</i> .	941
889		942
890	Jiho Shin, Hoeseok Yang, and Youngmin Yi. 2025a. GRASP: Group-based prediction of activation sparsity for fast LLM inference . In <i>Proceedings of the 62nd ACM/IEEE Design Automation Conference (DAC)</i> .	943
891		944
892		945
893		946
894		947
895	Jiho Shin, Hoeseok Yang, and Youngmin Yi. 2025b. Sparseinfer: Training-free prediction of activation sparsity for fast llm inference . In <i>Proceedings of the Design, Automation & Test in Europe Conference (DATE)</i> .	948
896		
897		
898		
899		
900	Susav Shrestha, Brian Settlemeyer, Joel Dapello, and Peng Chen. 2025. Polar sparsity: High throughput batched LLM inferencing with scalable contextual sparsity . In <i>Advances in Neural Information Processing Systems (NeurIPS)</i> .	949
901		950
902		951
903		952
904		
905	Chenyang Song, Xu Han, Zhengyan Zhang, Shengding Hu, Xiyu Shi, Kuai Li, Chen Chen, Zhiyuan Liu, Guangli Li, Tao Yang, and Maosong Sun. 2025. Prospare: Introducing and enhancing intrinsic activation sparsity within large language models . In <i>Proceedings of the International Conference on Computational Linguistics (COLING)</i> .	953
906		954
907		955
908		956
909		957
910		
911		
912	Yixin Song, Zeyu Mi, Haotong Xie, and Haibo Chen. 2024a. Powerinfer: Fast large language model serving with a consumer-grade gpu . In <i>Proceedings of the ACM Symposium on Operating Systems Principles (SOSP)</i> .	958
913		959
914		960
915		961
916		
917	Yixin Song, Haotong Xie, Zhengyan Zhang, Bo Wen, Li Ma, Zeyu Mi, and Haibo Chen. 2024b. Turbo sparse: Achieving llm sota performance with minimal activated parameters . <i>arXiv preprint arXiv:2406.05955</i> .	962
918		963
919		964
920		
921		
922	Mingjie Sun, Zhuang Liu, Anna Bair, and J. Zico Kolter. 2024. A simple and effective pruning approach for large language models . In <i>Proceedings of the International Conference on Learning Representations (ICLR)</i> .	965
923		966
924		967
925		968
926		969
927	Filip Szatkowski, Bartosz Wójcik, Mikołaj Piórczyński, and Simone Scardapane. 2024. Exploiting activation sparsity with dense to dynamic-k mixture-of-experts conversion . In <i>Advances in Neural Information Processing Systems (NeurIPS)</i> , volume 37.	970
928		971
929		972
930		
931		
932	Qwen Team. 2024. Qwen2.5 technical report . <i>arXiv preprint arXiv:2412.15115</i> .	973
933		974
934	Rongxiang Wang, Kangyuan Shu, and Felix Xiaozhu Lin. 2025. Enabling dynamic sparsity in quantized LLM inference . <i>arXiv preprint arXiv:2511.04477</i> .	975
935		976
936		977
	Xiuying Wei, Yunchen Zhang, Yuhang Li, Xiangguo Zhang, Ruihao Gong, Jinyang Guo, and Xianglong Liu. 2023. Outlier suppression+: Accurate quantization of large language models by equivalent and effective shifting and scaling . <i>arXiv preprint arXiv:2304.09145</i> .	978
		979
	Taiqiang Wu, Chaofan Tao, Jiahao Wang, Runming Yang, Zhe Zhao, and Ngai Wong. 2025a. Rethinking kullback-leibler divergence in knowledge distillation for large language models . In <i>Proceedings of the International Conference on Computational Linguistics (COLING)</i> .	980
		981
		982
	Tong Wu, Yutong He, Bin Wang, and Kun Yuan. 2025b. Mixture-of-channels: Exploiting sparse FFNs for efficient LLMs pre-training and inference . <i>arXiv preprint arXiv:2511.09323</i> .	983
		984
		985
		986
	Haotian Xu, Jiannan Yang, Tian Gao, Tsui-Wei Weng, and Tengfei Ma. 2025. Resting neurons, active insights: Robustify activation sparsity for large language models . In <i>Proceedings of the International Conference on Machine Learning (ICML)</i> .	987
		988
		989
		990
		991
		992
	Zhiyuan Xue, Yixin Song, Zeyu Mi, Le Chen, Yubin Xia, and Haibo Chen. 2024. Powerinfer-2: Fast large language model inference on a smartphone . <i>arXiv preprint arXiv:2406.06282</i> .	
	others Yang. 2026. WiSparse: Weight-aware mixed activation sparsity for large language models . <i>arXiv preprint arXiv:2602.14452</i> .	
	Lu Yin, You Wu, Zhenyu Zhang, Cheng-Yu Hsieh, Yaqing Wang, Yiling Jia, Gen Li, Ajay Jaiswal, Mykola Pechenizkiy, Yi Liang, Michael Bendersky, Zhangyang Wang, and Shiwei Liu. 2024. Outlier weighed layerwise sparsity (OWL): A missing secret sauce for pruning LLMs to high sparsity . In <i>Proceedings of the International Conference on Machine Learning (ICML)</i> .	
	Yilong Zhang, Junyuan Shang, Huanqi Cao, Shiyao Li, Shuguang Jiang, Lingxiao Ma, Shun Zheng, Cheng Li, and Jiang Bian. 2025a. COTS: Connected sparse activation for efficient LLM inference . <i>arXiv preprint arXiv:2502.14475</i> .	
	Zhenyu Zhang, Zechun Liu, Yuandong Tian, Harshit Khaitan, Zhangyang Wang, and Steven Li. 2025b. R-sparse: Rank-aware activation sparsity for efficient llm inference . In <i>Proceedings of the International Conference on Learning Representations (ICLR)</i> .	
	Haocheng Zheng and 1 others. 2026. ActTail: Heavy-tailed self-regularization for activation sparsity in large language models . <i>arXiv preprint arXiv:2603.12272</i> .	
	Yanqi Zhou, Tao Lei, Hanxiao Liu, Nan Du, Yanping Huang, Vincent Zhao, Andrew M. Dai, Zhifeng Chen, Quoc V. Le, and James Laudon. 2022. Mixture-of-experts with expert choice routing . In <i>Advances in Neural Information Processing Systems (NeurIPS)</i> , volume 35.	

A Inference Efficiency

Table 5 reports inference throughput on a single NVIDIA RTX PRO 6000 Blackwell Server Edition GPU (96 GB; batch size 1, sequence length 2048, fp16, averaged over 100 forward passes).

Table 5: Inference latency on LLaMA-3.1-8B (RTX PRO 6000 Blackwell, batch=1, seq=2048, fp16, 100 forward passes). Overhead is relative to dense inference.

Method	Sparsity	ms/fwd	tok/s	Overhead
Dense	0%	119.11±0.57	17,195	—
KL predictor-only	30%	126.64±0.70	16,172	+6.3%
KL predictor-only	50%	126.78±0.31	16,154	+6.4%
Predictor + compensation	30%	129.94±0.75	15,761	+9.1%
Predictor + compensation	50%	129.71±0.53	15,789	+8.9%

The full KL-sparse pipeline adds 6.4% overhead to dense inference (126.78 vs. 119.11 ms/forward), of which the predictor’s 32-layer forward pass contributes 3.17 ms (2.7%); the remainder arises from mask application and sparse FFN indexing. Eliminating the compensation network saves an additional ~2.7% latency overhead and 67.2M parameters (the per-layer GEMV required for output rescaling); output-distribution KL thus provides both quality and efficiency advantages over BCE-based approaches. The overhead is near-identical at 30% and 50% sparsity because the predictor runs the same computation regardless of target sparsity; only the top- k threshold changes.

Note that the current implementation uses unstructured element-wise masks, which do not translate to wall-clock speedup without sparse kernel support. The pipeline adds ~6.4% overhead; net acceleration requires kernel-level sparse computation optimizations (e.g., Triton block-sparse, cuSPARSE) that are beyond the scope of this work. Our contribution is prediction quality and inference simplification (removing the compensation network), not end-to-end latency reduction.

B Predictor Bottleneck Sensitivity

Table 6 reports the effect of predictor bottleneck dimension d_b on PPL at 50% sparsity under both TEAL global allocation (the paper’s primary inference protocol) and per-token evaluation.

Table 6: Predictor bottleneck sensitivity at 50% sparsity (WikiText-2 PPL↓). TEAL global is the primary inference protocol; per-token shown for reference.

d_b	Params	TEAL global	Δ	Per-token
64	38.2M (0.48%)	12.076	+0.9%	10.16
128	75.9M (0.95%)	11.963	—	9.218
256	151.5M (1.89%)	12.105	+1.2%	8.586
512	302.5M (3.77%)	12.315	+2.9%	7.93

Under TEAL global allocation, increasing d_b beyond 128 provides no benefit (PPL slightly worsens), while per-token evaluation shows monotonic improvement (10.16→7.93). This divergence indicates that the global top- k allocation mechanism, not predictor capacity, is the binding constraint: additional predictor capacity produces better per-neuron scores, but the global ranking and budget distribution cannot exploit them. The $d_b=128$ configuration (75.9M params, <1% of base model) is therefore optimal for the TEAL inference protocol.

B.1 Temperature Schedule Sensitivity

Table 7: Temperature schedule sensitivity at 50% sparsity (WikiText-2 PPL↓, TEAL global).

τ Schedule	PPL↓
Fixed $\tau=0.5$	11.861
Linear 1.0 → 0.1 (default)	11.963
Fixed $\tau=0.1$	12.072

Temperature schedule has minimal impact (<1.8% PPL variation). The linear annealing default is a conservative choice; even removing annealing entirely (fixed $\tau=0.5$) yields comparable performance.

B.2 Training Convergence

Table 8: Training convergence at 50% sparsity (WikiText-2 PPL↓, TEAL global).

Steps	PPL↓
200	15.158
400	12.538
600	12.046
800	11.907
1000	11.906

Training converges rapidly: 95% of the quality gap closes by step 600, and performance saturates by step 800 ($\Delta < 0.001$ between 800 and 1000 steps). Our default of 1000 steps provides a conservative margin.

C Geometric Penalty Schedule Dynamics

Table 9 reports the convergence of the geometric penalty schedule across sparsity targets. The geometric doubling pattern (1 → 2 → 4 → 8 → ...) enables λ to reach steady state within ~10 steps; the deadzone ($\epsilon = 0.02$) then stabilizes λ once actual sparsity is within $\pm 2\%$ of the target.

At 50% sparsity, the SwiGLU gate distribution already clusters near the target, so λ remains low (8); the constraint is nearly satisfied by the initial predictor outputs. At 30% and 70%, the gap between the gate prior and target is larger, driving λ

Table 9: Geometric penalty schedule convergence across sparsity targets (margin $\epsilon = 0.02$, 1000 training steps). Deviation is $|\text{actual} - \text{target}| / \text{target}$.

Target ρ	$\lambda_{\max}$	Final λ	Actual sparsity	Deviation
30%	5000	5000 (cap)	0.285	5.0%
40%	5000	625	0.398	0.5%
50%	1000	8	0.504	0.8%
70%	1000	1000 (cap)	0.688	1.7%

to its respective caps. The 40% target converges to $\lambda = 625$ with only 0.5% deviation; this precision reflects a regime where the target is reachable without λ saturation.




Figure 4: Lagrangian controller dynamics for 30% sparsity target ($\lambda_{\max} = 5000$). Left axis (log scale): penalty weight λ rises from 0.01 to its cap once actual sparsity undershoots the target at $\sim$ step 150, exhibiting bang-bang oscillations before saturating. Right axis: actual sparsity (solid) converges from the natural density (≈ 0.41) toward the target (dashed) and stabilizes near 0.28.

Temperature τ is annealed linearly from 1.0 to 0.1 over 1000 steps. A mild loss increase in the final ~ 100 steps is expected: as masks become more binary ($\tau \rightarrow 0$), discrete rounding noise slightly increases KL divergence, consistent with convergence toward the hard-mask inference regime.

The controller reliably drives actual sparsity to within 5% of the target across all tested regimes, with λ spanning three orders of magnitude (8 to 5000) via geometric doubling. For the 30% target, actual sparsity is $\approx 28.5\%$ (5% deviation). This does not affect conclusions: TEAL at 28.5% actual sparsity yields PPL 10.30 (vs. 10.97 at exact 30%), a 0.67 PPL difference negligible relative to the KL vs. TEAL gap (8.96 vs. 10.97).

D Out-of-Distribution Evaluation

To address concerns about data homogeneity between training (WikiText-103) and evaluation (WikiText-2), we evaluate on C4 (Raffel et al., 2020) validation set under TEAL global allocation. Table 10: Out-of-distribution evaluation on C4 validation (PPL $\downarrow$). 30% KL reports 2-seed mean $\pm$ std; other entries use seed 42.

Sparsity	Method	C4 PPL $\downarrow$
0%	Dense	9.21
30%	KL predictor (Ours)	16.03$\pm$0.01
50%	KL predictor (Ours)	18.40
	TEAL vanilla	38.75

At the 30% sparsity level, the KL predictor achieves C4 PPL 16.03 $\pm$ 0.01 (2-seed, global allocation, $1.74\times$ dense), confirming OOD generalization. At 50%, the KL predictor achieves 52.5% relative advantage over TEAL (18.40 vs. 38.75), comparable to the 50% advantage on WikiText-2 (11.92 vs. 23.90), confirming that KL-trained predictions generalize across distributions.

Uniform Allocation. Under uniform allocation at 30%, C4 PPL across three seeds is 18.22 (s42), 16.37 (s123), 16.36 (s456), yielding mean 16.98 ± 1.07 ($1.84\times$ dense). The cross-seed variance is substantially higher than in-domain: CV=6.3% on C4 vs. 0.2% on WikiText-2 (PPL 8.81 $\pm$ 0.02), driven by seed 42’s elevated C4 PPL (18.22 vs. 16.37 for other seeds). Global allocation’s cross-layer budget flexibility partially compensates for domain shift (PPL 16.03 global vs. 16.98 uniform), consistent with the Limitations observation that the predictor partially overfits to training-domain activation patterns.

Downstream Task Robustness. Table 11 compares downstream performance when the predictor generates masks from C4 text (OOD) vs. WikiText (in-domain), both at 30% uniform allocation (seed 42). All six benchmarks show $\Delta \leq 1.1\text{pp}$, confirming that predictor masks are robust to calibration domain shift.

E Per-Layer Retention Analysis

F Training Stability

Three-seed evaluation ($s \in \{42, 123, 456\}$) at 50% sparsity under uniform allocation yields PPL 11.92 ± 0.07 (CV=0.6%); training is stable across random initializations. Downstream metrics are similarly stable: ARC-c $33.8 \pm 1.2\%$, WinoGrande

Table 11: Downstream task robustness to calibration domain shift (KL 30% uniform, seed 42). WikiText: predictor masks from in-domain text; C4: masks from OOD text. Metrics: acc_norm (ARC-c, HS, PIQA), acc (WG, MMLU), exact_match (GSM8K).

Calibration	ARC-c	WG	MMLU	GSM8K	HS	PIQA
WikiText	40.8	68.6	40.6	7.4	70.9	74.8
C4 (OOD)	40.4	69.3	41.4	6.3	70.7	74.0
Δ	-0.4	+0.7	+0.8	-1.1	-0.2	-0.8




Figure 5: Per-layer neuron retention under global top- k allocation at 30%, 50%, and 70% sparsity. Dashed lines indicate uniform retention baselines. Late layers (L24–31) are consistently starved due to long-tailed activation distributions, with retention dropping to 28.4% (L31) even at 30% sparsity. At 70%, starvation extends to mid-layers (L10–12: 16.7–20.1%), explaining the transition to the capacity-limited regime.

62.9 $\pm$ 0.1%, MMLU 28.0 $\pm$ 0.6%, GSM8K 2.2 $\pm$ 0.1%.

Under global allocation at 50%, MMLU is 28.3 $\pm$ 0.7%; the near-random performance is capacity-limited rather than seed-sensitive (σ =0.7pp vs. 9.5pp at 30% global). The low standard deviation across both allocation strategies indicates that the KL training landscape is well-conditioned for this predictor architecture, with low sensitivity to initialization.

F.1 Cross-Seed Mask Consistency

To assess whether KL predictors discover consistent mask patterns across initializations, we compute pairwise Jaccard distances among three independently trained KL predictors (s42/s123/s456) at 30% uniform allocation (100 sequences, WikiText-2 validation).

The mean cross-seed Jaccard distance is 0.363 (IoU = 0.637), substantially lower than the KL-vs-magnitude divergence (0.46–0.68), confirming that KL-discovered patterns form a distinct cluster from magnitude-based masks. Per-layer divergence follows a U-shaped pattern: lowest at embedding-

adjacent layers (L0: 0.33), peaking at early layers (L1: 0.41), with a mid-network valley (L12–14: 0.33). The random baseline Jaccard distance is 0.462.

While specific mask configurations vary moderately across seeds (Jaccard distance 36.3%), the resulting PPL is highly stable (8.81 $\pm$ 0.02, CV=0.2%), suggesting multiple near-equivalent optima in the activation mask space. The cross-seed divergence (36.3%) is substantially lower than the KL-vs-magnitude divergence (46–68%), confirming that KL discovers effective patterns that cluster distinctly from magnitude-based selection.

G MSE Baseline Configuration

The MSE predictor uses identical architecture (two-layer MLP, bottleneck d_b =128), training data (WikiText-103, 1000 steps), optimizer (AdamW, lr=1 $\times$ 10⁻³, weight decay 1 $\times$ 10⁻², cosine decay), Gumbel-Sigmoid mask generation (τ : 1.0 $\rightarrow$ 0.1), and adaptive Lagrangian sparsity control as the KL and BCE predictors. The only difference is the loss target: per-layer $\|\mathbf{W}_d^{(\ell)}(\tilde{\mathbf{m}}^{(\ell)} \odot \mathbf{g}^{(\ell)} \odot \mathbf{x}^{(\ell)} \mathbf{W}_u^{(\ell)}) - \mathbf{y}_{\text{dense}}^{(\ell)}\|_2^2$, replacing the end-to-end KL divergence. This controlled comparison isolates the effect of output-distribution alignment: all three predictors see identical data in identical order with identical architectures; only the loss function differs.

MSE uses λ_{max} =1000 (vs. KL’s 5000) because MSE’s per-layer gradient magnitude is approximately 5 $\times$ larger than KL’s end-to-end gradient, requiring proportionally less penalty pressure to achieve comparable sparsity control. Both reach their target sparsity within 110 steps. At 50% target, MSE converges to the target sparsity for seed 42; however, with seed 123, sparsity reaches only 5.1% under the same λ_{max} cap (eval PPL=15,789, random-level), indicating high sensitivity to initialization—a failure mode not observed with KL-based losses. At 30%, MSE achieves 27.7% actual sparsity for s42 (gap 2.3pp). The 3-seed MSE mean (18.52 $\pm$ 1.68; s42=20.43, s123=17.73, s456=17.39) is still 1.7 $\times$ worse than TEAL (10.82), confirming that the performance deficit stems from the loss function, not insufficient sparsity pressure.

H Extended SNR Analysis

Figure 6 plots gradient SNR across the full 1000-step training for both KL and BCE predictors, measured at 14 checkpoints (LLaMA-3.1-8B, C4) with




Figure 6: Gradient SNR trajectory across 1000 training steps. BCE (orange) maintains consistently higher SNR than KL (blue) from step 100 onward ($2.3\text{--}17\times$), yet KL produces superior masks (Table 2), suggesting that target alignment matters more than gradient clarity for mask quality.

95% bootstrap confidence intervals. From step 100 onward, BCE maintains consistently higher SNR than KL ($2.3\text{--}17\times$ depending on training phase), with the gap narrowing toward convergence ($2.3\times$ at step 1000: 2.026×10^{-3} vs. 8.78×10^{-4}). Despite KL’s noisier gradient signal, KL-trained predictors produce substantially better masks (Table 2), confirming that *target alignment* matters more than gradient clarity for mask quality. BCE’s higher SNR reflects its simpler per-neuron classification objective, which concentrates gradients on individual neuron decisions; KL’s coupled gradient distributes signal across all mask decisions jointly, yielding lower SNR but discovering compositionally superior mask patterns (§3.6).

I Extended Downstream Evaluation

Table 12 reports HellaSwag and PIQA across sparsity levels and allocation strategies.

Table 12: Extended downstream benchmarks (LLaMA-3.1-8B). HellaSwag: 10-shot normalized acc; PIQA: 0-shot normalized acc. KL 30% uniform: 3-seed mean $\pm$ std; others: seed 42.

Sparsity	Method	HellaSwag	PIQA
0%	Dense	82.35	80.90
30%	KL uniform	71.07$\pm$0.54	74.95$\pm$0.25
	TEAL uniform	63.31	72.80
	BCE isocompute	62.81	74.21
50%	KL global	58.95	69.48
	KL uniform	56.66	69.10
	TEAL uniform	39.35	63.66

At 30%, the KL predictor retains 86.3% of dense HellaSwag accuracy (71.07 ± 0.54 vs. 82.35, 3-seed) and 92.6% of PIQA (74.95 ± 0.25 vs. 80.90),

confirming moderate degradation on non-reasoning tasks. At 50%, global allocation outperforms uniform on HellaSwag (+2.29pp), consistent with the budget-concentration benefit observed for PPL in the capacity-limited regime (§4.6); PIQA shows a smaller gap (+0.38pp).

J Compensation Training Convergence

All compensation networks (Gautam et al., 2026) were trained for 1000 steps ($\text{lr}=10^{-4}$, batch size 4, seq length 2048). Training loss plateaus by step 700 across all valid runs: 30% s42 (final loss 0.70), 30% s456 (0.68), 50% s42 (1.07), with fluctuations remaining within 0.08 loss units for the final 300 steps. The compensation network converges but provides no quality improvement over the KL predictor alone, confirming that the redundancy result reflects predictor self-sufficiency rather than insufficient compensation training.

K Qwen-2.5-14B Detailed Results

Table 13: Qwen-2.5-14B at 30% nominal sparsity. Actual sparsity differs: KL 29.6% (3-seed), TEAL 29.3% (target 30%). **Bold**: best sparse per column per group. BCE collapses to 6.4% native sparsity (§3.6), so forced global top- k is used.

Method	PPL $\downarrow$	ARC-c	WG	MMLU	GSM8K	HS	PIQA
Dense	5.19	67.58	81.06	79.77	87.57	84.20	81.94
<i>Global top-k allocation</i>							
KL (Ours)	7.62	57.17	72.06	63.72	48.82	75.01	77.37
TEAL	8.73	56.83	72.69	71.83	49.51	70.51	74.70
BCE	8.76	57.59	72.61	71.67	53.07	53.43	68.28
<i>Uniform allocation</i>							
KL (Ours)	7.90	53.50	72.82	63.55	32.2 $\pm$ 5.3	72.68$\pm$0.2	76.01$\pm$1.1
TEAL	8.01	55.29	69.85	66.34	52.1	67.72	74.70

GSM8K: 3-seed (s42/s123(v4)/s456).

Under global allocation, BCE’s HellaSwag drops to 53.43% (−17pp vs. TEAL), indicating that the magnitude-based predictor’s mask quality degrades severely when combined with non-uniform budget distribution at 14B scale.

L Full Allocation Analysis

Table 15 extends Table 14 with all downstream metrics.

M Mask Pattern Visualization

Figure 8 shows per-layer Jaccard distance between KL-derived and magnitude-based masks at 30% and 50% sparsity. Figure 7 shows per-neuron importance scores at layer 16. Figure 9 shows binary mask comparisons at layers 2, 16, and 28.

Table 14: Allocation summary (LLaMA-3.1-8B). $\pm$: 3-seed; *high seed variance.

Sparsity	Allocation	PPL $\downarrow$	MMLU
30%	Global	8.96 ± 0.02	$35.7 \pm 9.5^*$
	Uniform	8.81 ± 0.02	40.7 ± 1.3
50%	Global	11.84 ± 0.15	28.3 ± 0.7
	Uniform	11.92 ± 0.07	28.0 ± 0.6




Figure 7: KL vs. magnitude importance scores for layer 16 neurons at 50% sparsity. 41% of neurons fall in disagreement quadrants where one method selects a neuron the other discards.

N Cross-Architecture Comparison

Table 16 evaluates on Mistral-7B-v0.3 (Jiang et al., 2023) (GQA + sliding window, dense PPL 5.52).

The KL PPL advantage ratio is consistent across architectures (within 8% at $\leq 50\%$ sparsity). Table 17 reports downstream task accuracy at 30% sparsity on Mistral-7B-v0.3 (3-seed). KL improves PPL (-17.3%) and commonsense tasks (WG $+6.7\text{pp}$) but MMLU is tied (43.3% vs. 43.5%); ARC-c shows high seed variance ($\sigma=6.4\text{pp}$); GSM8K reverses.

O Taxonomy Discussion

Correction granularity. *Static* methods (Xu et al., 2025) learn a fixed bias vector per layer that does not depend on the input token. *Per-token* methods (Liu et al., 2023; Gautam et al., 2026; Akhauri et al., 2024; Szatkowski et al., 2024) produce a unique mask for each input position, conditioning on the token’s hidden state at each layer.

Training signal coupling. *Decomposed* objectives (BCE) treat each neuron independently, optimizing per-element classification against an oracle mask. *Intermediate* objectives (MSE on hidden states) capture local coupling within a single layer but not end-to-end output effects. *Coupled* objec-

tives (KL on output distribution) propagate gradients through the full forward pass; each mask decision’s gradient reflects its joint effect with all other masks on the final distribution.

SPON (Xu et al., 2025) distills a per-layer bias via KL but cannot adapt to per-input neuron importance (static $\times$ coupled). DeJaVu (Liu et al., 2023) and FastForward (Gautam et al., 2026) achieve per-token granularity but use BCE, where each neuron’s gradient is independent of other mask states. D2DMoE (Szatkowski et al., 2024) provides intermediate coupling (MSE on hidden states) but not end-to-end output effects. We define *output-distribution alignment* as whether the training loss optimizes a quantity that composes across layers. KL on output logits directly optimizes next-token prediction loss, which does not suffer from error composition; per-layer MSE minimizes reconstruction error independently at each layer, but layer-wise optima do not compose—a mask minimizing FFN output error at layer ℓ may amplify errors in layers $\ell+1 \dots L$. The SNR analysis (§3.6) corroborates this: KL achieves the best masks despite $2.3\times$ lower SNR than BCE, while MSE produces the worst masks. The framework predicts that losses operating on output-level quantities should outperform intermediate-level losses, a hypothesis consistent with our MSE results and testable with alternative output-level objectives.

P Oracle Mask Provenance

C1’s oracle masks are extracted from a converged KL predictor and binarized via hard threshold (logits > 0 , equivalent to $\sigma > 0.5$); they reflect one specific training trajectory (optimizer state, data order, Gumbel noise). The s42/s123/s456 consistency ($\leq 1.70\text{pp}$) provides partial evidence of robustness, but does not rule out trajectory-dependent confounds.

Q SPON Orthogonality

SPON (Xu et al., 2025) bias correction is weakly correlated with the KL predictor (cosine similarity 0.196, $R^2=0.038$) and yields only 2.4% PPL improvement (12.08 $\rightarrow$ 11.79); the predictor already captures most recoverable signal. The decomposition confirms that SPON and the KL predictor operate on largely orthogonal error sources.

Table 15: Full allocation analysis: global top- k vs. uniform across sparsity levels (LLaMA-3.1-8B). Entries with $\pm$: 3-seed mean $\pm$ std ($s \in \{42, 123, 456\}$); others: seed 42. *High MMLU seed variance (individual seeds span 25.3–43.5%).

Sparsity	Allocation	PPL $\downarrow$	ARC-c	WG	MMLU	GSM8K
25%	Global	8.17	45.4	70.6	46.6	11.5
	Uniform	8.23	46.5	71.8	44.0	11.0
30%	Global	8.96 $\pm$ 0.02	43.8$\pm$3.6	70.2$\pm$0.9	35.7 $\pm$ 9.5*	8.0$\pm$0.7
	Uniform	8.81$\pm$0.02	43.7 $\pm$ 2.5	68.9 $\pm$ 0.4	40.7$\pm$1.3	7.5 $\pm$ 0.1
35%	Global	9.31	38.9	67.6	40.4	7.7
	Uniform	9.44	38.7	66.5	37.8	4.3
40%	Global	10.14	39.5	67.5	39.5	6.6
	Uniform	10.76	39.6	63.1	35.5	3.6
45%	Global	10.80	39.2	64.6	30.4	3.0
	Uniform	10.95	37.6	63.6	30.6	2.7
50%	Global	11.84$\pm$0.15	36.3$\pm$3.2	63.8$\pm$0.3	28.3$\pm$0.7	2.7$\pm$1.3
	Uniform	11.92 $\pm$ 0.07	33.8 $\pm$ 1.2	62.9 $\pm$ 0.1	28.0 $\pm$ 0.6	2.2 $\pm$ 0.1




Figure 8: Per-layer Jaccard distance between KL-derived and magnitude-based masks at 30% and 50% sparsity (64 samples, 131K tokens). Divergence peaks in layers 24–28 (Jaccard distance >0.73 at 50%) and increases with sparsity across all layers.

Table 16: Cross-architecture comparison (WikiText-2 PPL $\downarrow$). Adv = TEAL / KL ratio. **Note:** LLaMA results use global top- k allocation; Mistral results use uniform allocation. Mistral 30%: 3-seed mean $\pm$ std; other rows: seed 42.

	Mistral-7B			LLaMA-8B		
	KL	TEAL	Adv	KL	TEAL	Adv
30%	6.93$\pm$0.19	8.38	1.21 $\times$	8.96	10.97	1.22 $\times$
50%	8.65	18.31	2.12 $\times$	11.79	23.31	1.98 $\times$
70%	13.33	61.72	4.63 $\times$	17.91	100.11	5.59 $\times$

R Per-Layer KL Downstream Decomposition

Table 18 reports raw downstream numbers for the signal type vs. coupling scope decomposition at 30% sparsity (§4.5). BoolQ’s $\approx 100\%$ coupling is an upper bound confounded by high cross-seed variance ($\pm 11.5\text{pp}$ for per-layer KL); the $K=4$ pilot (Table 19) suggests much of this gap is train–inference mismatch.

Table 17: Mistral-7B-v0.3 downstream evaluation at 30% sparsity. $\pm$: 3-seed mean $\pm$ std ($s \in \{42, 123, 456\}$); TEAL is deterministic (single run). ARC-c exhibits high seed variance ($\sigma=6.4\text{pp}$).

Method	PPL $\downarrow$	ARC-c	WG	HS	PIQA	MMLU	GSM8K
Dense	5.52	60.2	78.7	83.2	81.9	62.4	36.5
KL 30%	6.93$\pm$0.19	44.9$\pm$6.4	70.2$\pm$1.2	69.7$\pm$0.8	74.8$\pm$0.4	43.3 $\pm$ 1.5	8.6 $\pm$ 2.2
TEAL 30%	8.38	41.1	63.5	60.1	72.1	43.5	15.9

S $K=4$ Mismatch Control

Table 19 reports 3-seed $K=4$ per-layer KL results at 30% sparsity. Training with $K=4$ random layers per step reduces the train–inference mismatch (4/32 vs. 1/32 layers masked during training, while inference applies all 32). The gap closure column reports the fraction of the $K=1 \rightarrow \text{E2E}$ KL gap closed by $K=4$.

T Per-Layer Sparsity Constraint Analysis

To test the causal role of the KL predictor’s learned per-layer allocation, we evaluate two inference-time interventions at 30% overall sparsity (seed 42, uniform allocation baseline):




Figure 9: Binary mask comparison between KL-derived and magnitude-based masks at selected layers (L2, L16, L28) for 30% and 50% sparsity (64 samples). Black = selected by both, light regions = selected by only one method. Deep layers show the largest divergence (L28 Jaccard distance: 0.60 at 30%, 0.74 at 50%).

Table 18: 2D decomposition at 30% sparsity (LLaMA-3.1-8B, uniform allocation). All rows are 3-seed mean $\pm$ std. Signal type = per-layer KL – MSE; coupling = E2E KL – per-layer KL. Percentages are upper bounds (dimensions not fully orthogonal; see §4.5).

	PPL↓	ARC-c	WG	MMLU	GSM8K	BoolQ
E2E KL	8.81 $\pm$ 0.02	43.7	68.9	40.7	7.5	73.7
Per-layer KL	9.05 $\pm$ 0.02	39.9	65.9	38.6	5.8	64.5
MSE	18.52 $\pm$ 1.68	35.8 $\pm$ 4.1	59.8 $\pm$ 2.6	30.7 $\pm$ 2.6	2.3 $\pm$ 0.7	64.5 $\pm$ 5.3
Total gap	9.71	7.9	9.1	10.0	5.2	9.2
Signal type (%)	97.5	51.9	67.0	79.0	67.3	$\approx$ 0
Coupling (%)	2.5	48.1	33.0	21.0	32.7	$\approx$ 100

Protecting late layers (L22–L28) does not improve knowledge-intensive metrics despite reducing their sparsity from 48–57% to $\leq 30\%$, while worsening PPL by 5.2%. Conversely, removing the protection valley (forcing L14–L18 to 30% instead of 14–18%) improves MMLU by +3.4pp and GSM8K by +1.6pp at the cost of ARC-c – 1.5pp. The KL predictor’s learned allocation thus represents a task-dependent trade-off—prioritizing commonsense reasoning via middle-layer protection—rather than a uniformly optimal pattern. This allocation trade-off also explains the $K=4$ MMLU anomaly (§S): $K=4$ ’s near-uniform learned sparsity (std=0.003) inherently removes the valley, producing MMLU=43.5 comparable to the remove-valley condition (44.0).

U Per-Layer Sparsity Distribution

The KL predictor at 30% target sparsity learns a non-uniform per-layer allocation that differs markedly from magnitude-based selection. Aggre-

Table 19: 3-seed $K=4$ mismatch control at 30% sparsity (LLaMA-3.1-8B, uniform allocation, $s \in \{42, 123, 456\}$). Actual sparsity 29.6%; per-layer sparsity std=0.003 (near-uniform). Gap closure = $(K=4 - K=1)/(E2E - K=1) \times 100$.

	$K=4$ (3-seed)	$K=1$ (3-seed)	E2E KL (3-seed)	Gap closure (%)
PPL↓	8.89 $\pm$ 0.01	9.05 $\pm$ 0.02	8.81 $\pm$ 0.02	66.7
ARC-c	40.0 $\pm$ 1.3	39.9	43.7 $\pm$ 2.5	2.6
WG	66.9 $\pm$ 0.2	65.9	68.9 $\pm$ 0.4	33.3
MMLU*	42.7 $\pm$ 1.4	38.6	40.7 $\pm$ 1.3	>100
GSM8K	7.0 $\pm$ 1.6	5.8	7.5 $\pm$ 0.1	70.6
BoolQ†	67.7 $\pm$ 7.5	64.5 $\pm$ 11.5	73.7	34.8

*MMLU $K=4$ systematically exceeds E2E across all 3 seeds;

attributable to $K=4$ ’s near-uniform per-layer sparsity (std=0.003) matching uniform eval allocation. Independently confirmed: removing the predictor’s middle-layer valley yields MMLU=44.0% (Appendix T). †BoolQ exhibits high inter-seed variance for both $K=4$ (std=7.5) and $K=1$ (std=11.5); gap closure unreliable.

gate sparsity is similar across broad regions (early L0–L7: KL 28.1% vs. magnitude 27.5%; late L24–L31: 44.9% vs. 47.2%), but middle layers (L8–L23) exhibit high internal variance despite similar averages (28.2% vs. 27.3%). Table 21 highlights three sub-ranges.

V C4 Training Data Control

To test whether WikiText-103’s narrow domain drives the knowledge-task regression (§4), we train the KL predictor on C4 (Raffel et al., 2020) (diverse web text) under identical hyperparameters (1000 steps, 30% uniform, s42).

The C4-trained predictor achieves comparable

Table 20: Per-layer sparsity constraint analysis. *Protect late*: cap L22–L28 sparsity $\leq 30\%$; *Remove valley*: override L14–L18 sparsity to 30%. Budget redistributed to maintain 30% overall.

Condition	PPL	ARC-c	MMLU	GSM8K
KL baseline (s42)	8.82	40.8	40.6	7.4
Protect late layers	9.26	42.4	40.7	6.5
Remove valley	8.76	39.3	44.0	9.0

Table 21: Per-layer sparsity sub-ranges: KL predictor vs. magnitude (LLaMA-3.1-8B, 30%, uniform allocation). Ranges denote min–max across layers in each group.

Layer Range	KL (%)	Mag (%)	Δ (pp)
L8–L12 (high pruning)	32–36	25–26	+7 to +10
L14–L18 (protection valley)	14–18	28–30	–8 to –16
L22–L28 (aggressive pruning)	48–57	—	—

MMLU (40.0% vs. 40.6%) but worse GSM8K (4.1% vs. 7.4%) and higher PPL (11.55 vs. 8.82), despite training on more diverse text. This rules out training data composition as the primary cause of knowledge-task regression and supports the hypothesis that forward KL’s mode-covering property structurally under-protects rare reasoning-critical activations.

W Per-Layer Coupling Breadth Analysis

We vary the number of layers coupled per training step (K) from 1 (pure per-layer) to 32 (end-to-end), evaluating all predictors under TEAL uniform allocation at 30% sparsity on LLaMA-3.1-8B.

PPL exhibits a non-monotonic U-shape (Figure 2): $K=2$ is worst (10.75 ± 0.15), likely due to insufficient per-layer updates (~ 62 updates/layer vs. ~ 250 at $K=8$). $K=16$ achieves the lowest PPL (8.67 ± 0.02), marginally better than E2E (8.81 ± 0.02). $K=16$ MMLU exhibits high seed variance (std=2.9; s42=44.0, s123=40.0, s456=38.7), with the 3-seed mean (40.9) indistinguishable from E2E (40.7 ± 1.3). The signal type gap ($K=1$ KL 9.05 vs. per-layer MSE 18.52 ± 1.68) of 9.47 PPL exceeds the coupling breadth range within KL (best $K=16$ 8.67 vs. worst $K=2$ 10.75 = 2.08 PPL), reinforcing that output-distribution alignment, not coupling scope, is the primary quality lever for PPL at this sparsity level.

X Licenses

LLaMA-3.1-8B is released under the Llama 3.1 Community License Agreement. Mistral-7B-v0.3 is released under the Apache 2.0 License. Qwen-2.5-14B is released under the Apache 2.0 License. WikiText-103 and WikiText-2 (Merity et al., 2017) are released under the Creative Commons

Table 22: C4-trained vs. WikiText-trained KL predictor (LLaMA-3.1-8B, 30% uniform, s42).

Training Data	PPL↓	ARC-c	WG	MMLU	GSM8K	HS / PIQA
WikiText-103	8.82	40.78	68.59	40.6	7.35	70.85 / 74.81
C4	11.55	43.09	68.67	39.95	4.09	69.24 / 75.03

Table 23: K-curve: coupling breadth vs. predictor quality (30%, TEAL uniform, LLaMA-3.1-8B). All values are 3-seed mean $\pm$ std (s42/s123/s456) except $K=1$ downstream (single seed). “—” = not evaluated.

K	PPL↓	ARC-c	WG	MMLU	GSM8K	HS	PIQA
1	9.05 ± 0.02	39.9	65.9	38.6	5.8	—	—
2	10.75 ± 0.15	39.7 ± 1.2	63.2 ± 1.7	38.3 ± 3.2	2.8 ± 0.6	66.5 ± 1.0	73.5 ± 0.7
4	8.89 ± 0.01	40.0 ± 1.2	66.9 ± 0.2	42.7 ± 1.4	7.0 ± 1.6	67.8 ± 1.0	74.5 ± 0.5
8	9.00 ± 0.06	40.7 ± 1.0	68.7 ± 1.3	40.4 ± 0.6	6.8 ± 1.3	69.0 ± 1.1	75.1 ± 1.1
16	8.67 ± 0.02	41.9 ± 1.1	68.5 ± 0.9	40.9 ± 2.9	8.5 ± 1.2	70.1 ± 0.4	75.0 ± 0.8
32 (E2E)	8.81 ± 0.02	43.7 ± 2.5	68.9 ± 0.4	40.7 ± 1.3	7.5 ± 0.1	—	—

Attribution-ShareAlike License (CC BY-SA 3.0). All evaluation benchmarks (ARC, WinoGrande, MMLU, GSM8K, HellaSwag, PIQA) are publicly available under their respective licenses. Our code will be released under the MIT License upon acceptance.

Y Use of AI Assistants

We used Claude (Anthropic) as an AI coding and writing assistant throughout this work. Specifically:

- **Code:** drafting experiment scripts, data-processing pipelines, and analysis utilities; debugging and refactoring.
- **Writing:** drafting and revising paper sections, polishing prose, and reformatting LaTeX.
- **Literature:** summarizing related papers to inform related-work section.